Yale

GLBL 7315 LAW 21473-01

The Global Law & Policy of Artificial Intelligence

Dr. Artur Pericles L. Monteiro
artur.monteiro@yale.edu

Spring 2024. *Last updated:* April 4, 2024 (3:51 pm).
Tuesday, 1:30–3:20. Watson Center A74.

ABOUT THIS COURSE

How does law shape algorithms, data, and artificial intelligence? And how is the law in turn transformed by AI? We are often told that law is too slow and fails to keep up with technology; that framing obscures such questions. Understanding the entanglements between law and technology and recognizing the paths that got us here offer valuable insights to different nations and the international community seeking to govern AI. This seminar course takes up those questions against the backdrop of geopolitical tension, as well as regulatory competition and imbalance.

COURSE OUTCOMES

- Develop a better understanding of how law has shaped and has been shaped by artificial intelligence and data.
- Analyze legal developments that played an important role in establishing conditions for the development and deployment of AI systems.
- Recognize how geopolitical dynamics are setting and constraining the possibilities each nation and the international community have to govern AI.
- Forge knowledge about current proposals for the governance and regulation artificial intelligence.

ENROLLMENT INFORMATION

Please share a short statement (2 or 3 sentences will work) about what interests you in the course and how you hope it will benefit your research agenda or professional goals. A resume or CV is welcome, but not required. I will use these materials to create a community that makes the most of the experience, perspective, and background of each member.

NOTE FOR JACKSON SCHOOL AND OTHER GRADUATE STUDENTS

This course will delve deep into the law. It is cross-listed with YLS and aims to bring together Jackson and the law school. We will be reading case opinions, law review articles, and memoranda. Rest assured that no legal training is assumed, and you will not be required to work on assignments that law students might work on. The goal is to afford you a first-hand grip on what some of the most important legal and policy issues are and how they are framed and disputed. An overarching argument in this course is that part of the technology we have before us, particularly under the umbrella of AI, can only be understood if we appreciate the shaping force of the law. This exercise should yield a greater understanding of the paths and prospects of AI, whatever your interests are.

All graduate students are welcome.

NOTE FOR LAW SCHOOL

The course has been designed to include both YLS and Jackson. It is a Jackson graduate course cross-listed with YLS, for 2 units.

YLS's calendar does not fully match Jackson's. The course is planned to accommodate that. The full class will meet together for 12 weeks, respecting YLS's reading period (April 23–29). We will discuss alternative activities for the week of March 10–17, when most of Yale will be on spring recess, though not YLS. See calendar.

J.D. and all graduate students are welcome.

NOTE FOR YALE COLLEGE

Advanced undergraduates will be considered after priority has been given to graduate and professional school students.

ATTENDANCE, ASSESSMENTS AND GRADING

Attendance is required and active participation is expected to create a meaningful community around the course. Notify me as soon as possible if something will prevent you from attending our sessions.

Assessments and grading are structured to reflect the centrality of community and engagement.

- *Short responses* (30%): Up to three short responses (500-800 words) reflecting on the readings will be assigned among class members to be submitted by Friday, 5 p.m., before the corresponding session. These will be used to guide the discussion in class.
- *Participation* (30%): Everyone is expected to engage with class discussions constructively, in dialogue with points and questions raised by other class members. Those assigned responses for that session will have a role in guiding class discussion.
- *Final paper* (40%): You will be free to choose between two options: a traditional essay and a policy memo. We will discuss these in class, including expectations and rubrics.
  - Traditional essay: A final paper critically reflecting on course discussions. 3500-5000 words. Longer papers are welcome if you want to meet extra credit requirements applying to your program.
  - Policy memo: A short (under 1500 words) brief or memo that effectively distills information or articulates strategy speaking, addressing a policymaker or other stakeholder. While shorter, this option entails research on the views, positions, and strategies of the intended audience.

COURSE RESOURCES

All readings will be available on Canvas. The following books are valuable resources for background and further reading.

- Julie E. Cohen, Between Truth and Power (2019).
- Solon Barocas, Moritz Hardt, & Arvind Narayanan, Fairness and Machine Learning: Limitations and Opportunities (2023).
- Kate Crawford, Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence (2021).
- Chris Wiggins, & Matthew L. Jones, How Data Happened: A History from the Age of Reason to the Age of Algorithms (2023).
- Louise Amoore, Cloud Ethics: Algorithms and the Attributes of Ourselves and Others (2020).
- Ifeoma Ajunwa, The Quantified Worker (2023).

CALENDAR (TENTATIVE)

		Jackson	YLS
Jan. 16	welcome		
Jan. 23	session I		
Jan. 30	session II		
Feb. 6	session III		
Feb. 13	session IV		
Feb. 20	session V		
Feb. 27	session VI		
Mar. 5	session VII		
Mar. 12		spring recess	final paper projects
Mar. 19		spring recess	spring recess
Mar. 26	session VIII		
Apr. 2	session IX		
Apr. 9	session X		
Apr. 16	session XI		
Apr. 23	policy memo session		YLS reading period
Apr. 30	policy memo session		YLS reading period
May 7	final paper		

PROGRAM

Note on readings While there is much more we could cover each week, I have made an effort to keep reading load reasonable. Readings for each session are around 50 pages. For a few sessions below with readings that go over 50 pages, I will still select excerpts from the texts to keep to that ballpark. The goal with this is making sure everyone can fully engage with the assigned readings each week.

Everyone should have read the *assigned reading*; *further reading* offers material that explores questions in more detail, provides different perspectives, or covers additional issues that are related to the session.

Readings are subject to change before the start of the course.

WELCOME & CONVERSATION

Our first session will be a conversation about the class's interests and expectations. A premise behind this course is that the study of artificial intelligence and its impacts on society must be multidisciplinary and seek understanding through a combination of knowledge forms and experiences. This applies to the class!

The first session will be an opportunity for class members to contribute to the syllabus drawing on their experience and expertise. This syllabus offers one approach to investigating AI governance and focuses on questions that I hope provide a good vantage point to consider the role of law has played and the role it can play. There is room for adaptation. We will discuss what topics the class is most interested in and what other themes could be incorporated along the semester. I will also invite suggestions for scholars who we might host to speak with the class.

In-class reading

- Julie E. Cohen, *Surveillance Capitalism as Legal Entrepreneurship*, 17 SURVEILLANCE & SOC'Y 240 (2019).

Further reading (optional)

- Jathan Sadowski, Salomé Viljoen, & Meredith Whittaker, *Everyone Should Decide How Their Digital Data Are Used — Not Just Tech Companies*, 595 NATURE 169 (2021).

I. THE LAW THAT MADE AI

Contrary to what is usually assumed, law does not always lag behind technology. In fact, accounting for the role of the law and appreciating how it changed in response to technologies is essential to understand how we got here. It will also be instrumental later when we consider governance frameworks.

Assigned reading

- Anupam Chander, *How Law Made Silicon Valley*, 63 EMORY L.J. 639, 641–669 (2014).
- Julie E. Cohen, BETWEEN TRUTH AND POWER 48–74 (2019).
- hiQ Labs, Inc. v. LinkedIn Corp., No 17-16783 (9th Cir. Sept. 9, 2020) (edited by James Grimmelmann).

Further reading

- Sheila Jasanoff, *Making Order: Law and Science in Action*, *in* HANDBOOK OF SCIENCE AND TECHNOLOGY STUDIES 761 (Edward J. Hackett et al. eds., 2007).
- Amy Kapczynski, *The Law of Informational Capitalism*, 129 YALE L.J. 1460, 1496–1514 (2019).
- Anupam Chander, *Section 230 and the International Law of Facebook*, 24 YALE J. L. & TECH. 393, 402–404, 416–420 (2022).
- Thomas E. Kadri, *Platforms as Blackacres*, 68 UCLA L. REV. 1184, 1194–1222 (2022).
- Salomé Viljoen, *The Promise and Limits of Lawfulness: Inequality, Law, and the Techlash*, 2 J. SOC. COMPUTING 284 (2021).
- Katharina Pistor, *Empire of Law*, *in* THE CODE OF CAPITAL: HOW THE LAW CREATES WEALTH AND INEQUALITY 1 (2019).

II. "DATA", DATAFICATION

Data fuels the current "AI spring." "Data-driven" decisions are not just seen as more rational, but increasingly as the only rational option. Datafication is an injunction defining of our time. Data is often presented as an objective representation, a reflection of society. This session probes such views and takes up the implications for claims about AI systems.

Assigned reading

- Lisa Gitelman, & Virginia Jackson, *Introduction*, *in* "RAW DATA" IS AN OXYMORON 1, 1–4 (Lisa Gitelman ed., 2013).
- Solon Barocas, & Andrew D Selbst, *Big Data's Disparate Impact*, 104 CALIF. L. REV. 671, 677–692 (2016).
- Jathan Sadowski, *When Data Is Capital: Datafication, Accumulation, and Extraction*, 6 BIG DATA & SOC'Y 1 (2019).
- Paul Dumouchel, *Data-Driven Agency and Knowledge*, *in* LIFE AND THE LAW IN THE ERA OF DATA-DRIVEN AGENCY 45 (Mireille Hildebrandt & Kieron O'Hara eds., 2020).

Further reading

- Solon Barocas, Moritz Hardt, & Arvind Narayanan, FAIRNESS AND MACHINE LEARNING: LIMITATIONS AND OPPORTUNITIES (2023), ch. 8.
- danah boyd, & Kate Crawford, *Critical Questions for Big Data: Provocations for a Cultural, Technological, and Scholarly Phenomenon*, 15 INFO. COMM. SOC'Y 662 (2012).
- Marion Fourcade, & Kieran Healy, *Seeing Like a Market*, 15 SOCIO-ECONOMIC REV. 9 (2016).
- Chris Wiggins, & Matthew L. Jones, *The Science of Data*, *in* HOW DATA HAPPENED 196 (2023).
- Aaron Mendon-Plasek, *Mechanized Significance and Machine Learning: Why It Became Thinkable and Preferable to Teach Machines to Judge the World*, *in* THE CULTURAL LIFE OF MACHINE LEARNING 31 (Jonathan Roberge & Michael Castelle eds., 2021).

III. GOVERNANCE OF WHAT? GOVERNANCE HOW?

We can think of governance in at least two ways — governance by algorithms and governance of algorithms. This session explores both. It starts our conversation about what governing algorithms can be.

Assigned reading

- Michael Veale, Kira Matus, & Robert Gorwa, *AI and Global Governance: Modalities, Rationales, Tensions*, 19 ANN. REV. L. & SOC. SCI., 1–11 (2023).
- Frank Pasquale, *The Second Wave of Algorithmic Accountability*, (2019), https://lpeproject.org/blog/the-second-wave-of-algorithmic-accountability/.
- Jack Balkin, *How to Regulate (and Not Regulate) Social Media*, 1 J. FREE SPEECH L. 71, 71–80 (2021).
- Alicia Solow-Niederman, *Administering Artificial Intelligence*, 93 S. CAL. L. REV. 633, 680–695 (2020).
- *Twitter, Inc. v. Taamneh*, 598 U.S. 471 (2023) (excerpts).

Further reading

- Solon Barocas, Moritz Hardt, & Arvind Narayanan, FAIRNESS AND MACHINE LEARNING: LIMITATIONS AND OPPORTUNITIES (2023), ch. 1.
- Helen Nissenbaum, *Accountability in a Computerized Society*, 2 SCI. & ENG'G ETHICS 25 (1996).
- Julia Angwin, et al. *Machine Bias*, PROPUBLICA (2016), https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing.
- Hannah Bloch-Wehba, *Transparency's AI Problem*, (2021).
- Ngozi Okidegbe, *Of Afrofuturism, of Algorithms*, 9 CRITICAL ANALYSIS L. 35 (2022).
- David Freeman Engstrom, & Amit Haim, *Regulating Government AI and the Challenge of Sociotechnical Design*, 19 ANN. REV. L. & SOC. SCI. 277 (2023).

IV. GOVERNANCE BY WHOM? A GLOBAL OVERVIEW

Current discussions about governing AI do not take place in a vacuum. The competition for who gets to regulate global technology is well underway. This session considers the ascendancy of the European Union and takes stock of the dominance so long exerted by the U.S. It also gives us our first peak into the EU's AI Act and reflects on the geopolitical rhetoric on China that has been driving pushes for action in the West. What does this mean for the Rest of the World?

Assigned reading

- Michael Veale, Kira Matus, & Robert Gorwa, *AI and Global Governance: Modalities, Rationales, Tensions*, 19 ANN. REV. L. & SOC. SCI., 11–17 (2023).
- Anu Bradford, *After the Fall of the American Digital Empire*, (2023), https://knightcolumbia.org/content/after-the-fall-of-the-american-digital-empire.
- Michael Veale, & Frederik Zuiderveen Borgesius, *Demystifying the Draft EU Artificial Intelligence Act — Analysing the Good, the Bad, and the Unclear Elements of the Proposed Approach*, 22 COMPUT. L. REV. INT'L 97 (2021).
- Jing Cheng, & Jinghan Zeng, *Shaping AI's Future? China in Global AI Governance*, 32 J. CONTEMP. CHINA 794 (2023).

Further reading

- Julian Posada, *The Coloniality of Data Work: Power and Inequality in Outsourced Data Production for Machine Learning*, Nov., 2022.
- European Data Protection Board, *Guidelines 3/2018 on the Territorial Scope of the GDPR (Article 3)*, (2019), https://edpb.europa.eu/sites/default/files/files/file1/edpb_guidelines_3_2018_territorial_scope_after_public_consultation_en_1.pdf.
- Marie-Therese Png, *At the Tensions of South and North: Critical Roles of Global South Stakeholders in AI Governance*, *in* THE OXFORD HANDBOOK OF AI GOVERNANCE (Justin B. Bullock et al. eds., 2024).
- Jonas Tallberg, et al. *The Global Governance of Artificial Intelligence: Next Steps for Empirical and Normative Research*, 25 INT'L STUD. REV. (2023).
- Robert Trager, et al. *International Governance of Civilian AI: A Jurisdictional Certification Approach*, SSRN ELECTRONIC JOURNAL (2023).

V. DATA PRIVACY: FUNDAMENTALS

Not so long ago, concerns about information technologies and their societal impacts were translated into a proto-governance of data. This session examines what those earlier pushes yield for data privacy (or data protection), in the U.S. and the EU It takes a closer look into the operation of the General Data Protection Regulation and discusses structural limitations to the approach it embodies.

Assigned reading

- Neil Richards, Andrew Serwin, & Tyler Blake, *Understanding American Privacy*, *in* RESEARCH HANDBOOK ON PRIVACY AND DATA PROTECTION LAW 60, 64–71 (Gloria González, Rosamunde Van Brakel, & Paul De Hert eds., 2022).
- Salomé Viljoen, *A Relational Theory of Data Governance*, 131 YALE L.J. 573, 577–586, 603–616 (2021).
- Anupam Chander, Margot E. Kaminski, & William McGeveran, *Catalyzing Privacy Law*, 105 MINN. L. REV. 1733, 1793–1802 (2021).
- European Data Protection Board, *Guidelines 05/2020 on Consent Under Regulation 2016/679*, 1–25 (2020), https://edpb.europa.eu/sites/default/files/files/file1/edpb_guidelines_202005_consent_en.pdf.

Further reading

- Court of Justice of the European Union, Judgment of 4 July 2023, *Meta Platforms, Inc. and Others v. Bundeskartellamt*, C-252/21, ECLI:EU:C:2023:537, §§ 86-126
- Sarah E. Igo, THE KNOWN CITIZEN 238–263 (2018).
- Daniel J Solove, *Privacy Self-Management and the Consent Dilemma*, 126 HARV. L. REV. 1880 (2013).
- Chris Wiggins, & Matthew L. Jones, HOW DATA HAPPENED: A HISTORY FROM THE AGE OF REASON TO THE AGE OF ALGORITHMS 141–174 (2023).

VI. DATA PRIVACY: MACHINE LEARNING

Data privacy is part of the existing law of artificial intelligence. This session considers how data privacy maps onto the issues that scholars have raised about artificial intelligence and big data.

Assigned reading

- Alicia Solow-Niederman, *Information Privacy and the Inference Economy*, 117 NW. U. L. REV., 378–400 (2022).
- Margot E. Kaminski, *The Right to Explanation, Explained*, 34 BERKELEY TECH. L.J. 189, 196–218 (2019).
- Bjørn Aslak Juliussen, Jon Petter Rui, & Dag Johansen, *Algorithms That Forget: Machine Unlearning and the Right to Erasure*, 51 COMPUT. L. & SEC. REV. (2023).

Further reading

- Solon Barocas, & Helen Nissenbaum, *Big Data's End Run Around Anonymity and Consent*, *in* PRIVACY, BIG DATA, AND THE PUBLIC GOOD: FRAMEWORKS FOR ENGAGEMENT 44 (2014).
- Sandra Wachter, Brent Mittelstadt, & Chris Russell, *Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR*, 31 841, 842–844, p.861–883 (2018).
- Marvin van Bekkum, & Frederik Zuiderveen Borgesius, *Using Sensitive Data to Prevent Discrimination by Artificial Intelligence: Does the GDPR Need a New Exception?*, 48 COMPUT. L. & SEC. REV. 105770, 5–12 (2023).
- Nadezhda Purtova, & Ronald Leenes, *Code as Personal Data: Implications for Data Protection Law and Regulation of Algorithms*, INT'L DATA PRIVACY L. (2023).
- Cornelius Witt, & Jan De Bruyne, *The Interplay Between Machine Learning and Data Minimization Under the GDPR: The Case of Google's Topics API*, INT'L DATA PRIVACY L. (2023).
- Salomé Viljoen, *A Relational Theory of Data Governance*, 131 YALE L.J. 573, 634–654 (2021).
- Daniel Susser, *From Procedural Rights to Political Economy: New Horizons for Regulating Online Privacy*, *in* THE ROUTLEDGE HANDBOOK OF PRIVACY AND SOCIAL MEDIA 281 (Sabine Trepte & Philipp Masur eds., 2023).
- Jamie Duncan, *Data Protection Beyond Data Rights: Governing Data Production Through Collective Intermediaries*, 12 INTERNET POL'Y REV., 9–17 (2023).
- Marvin van Bekkum, & Frederik Zuiderveen Borgesius, *Digital Welfare Fraud Detection and the Dutch SyRI Judgment*, 23 EUR. J. SOC. SEC. 323 (2021).

VII. GOVERNANCE HOW? FAIRNESS, TRANSPARENCY, AND ACCOUNTABILITY

Guest lecture: Professor Tim Rudner

Scholars who have been studying and reflecting on artificial intelligence have put forward frameworks to address the problems they identified. We turn our attention to the formulations they offered, which over time have deepened along the questions and scope of interventions they consider.

Assigned reading

- Joshua A Kroll, et al. *Accountable Algorithms*, 165 U. PA. L. REV. 633, 695–705 (2017).
- Mike Ananny, & Kate Crawford, *Seeing Without Knowing: Limitations of the Transparency Ideal and Its Application to Algorithmic Accountability*, 20 NEW MEDIA & SOC'Y 973 (2018).
- Tim Rudner, & Helen Toner, *Key Concepts in AI Safety: Interpretability in Machine Learning*, (2021), https://doi.org/10.51593/20190042.
- Finale Doshi-Velez, & Been Kim, *Towards a Rigorous Science of Interpretable Machine Learning*, (2017).
- Mike H. M. Teodorescu, & Christos Makridis, *Fairness in Machine Learning*, (2020), https://www.brookings.edu/articles/fairness-in-machine-learning-regulation-or-standards/.
- Inioluwa Deborah Raji, *The Anatomy of AI Audits: Form, Process, and Consequences*, *in* THE OXFORD HANDBOOK OF AI GOVERNANCE C28S1 (Justin B. Bullock et al. eds., 2024).
- Gianclaudio Malgieri, & Frank Pasquale, *Licensing High-Risk Artificial Intelligence: Toward Ex Ante Justification for a Disruptive Technology*, 52 COMPUT. L. & SEC. REV. 105899, 5–15 (2024).
- Neel Guha, et al. *The AI Regulatory Alignment Problem*, (2023), https://hai.stanford.edu/policy-brief-ai-regulatory-alignment-problem.

  ### Further reading

- Daniel Innerarity, *The Improvement of Democracy Through Transparency and Its Limits*, 5 INT'L J. PHIL. 44 (2017).

- Rebecca Williams, et al. *From Transparency to Accountability of Intelligent Systems: Moving Beyond Aspirations*, 4 Data & Pol'y (2022).
- Sonia K. Katyal, *Private Accountability in the Age of Artificial Intelligence*, 66 UCLA L. Rev. 56 (2019).
- Solon Barocas, Moritz Hardt, & Arvind Narayanan, Fairness and Machine Learning: Limitations and Opportunities 2 (2023).
- Emily Black, et al. *Less Discriminatory Algorithms*, (2024), https://dx.doi.org/10.2139/ssrn.4590481, manuscript.

VIII. GOVERNANCE HOW? RISKS AND SAFETY

Guest lecture: Professor Sofia Ranchordas

Risk regulation seems to lend itself well to the current discourse about AI. It is at the center of the EU's AI Act as well as other proposals that follow it. What does that modality of governance entail?

The proposed EU AI Act categorizes AI applications based on their potential risks. In this session, we will delve into this approach and focus on banned and "high-risk uses of AI." For example, the AI Act prohibits the use of AI for real-time biometric categorization, predictive policing based solely on profiling or personality traits, "except when supporting human assessment based on objective, verifiable facts linked to criminality," "facial recognition databases based on untargeted scraping." In the wake of several scandals in the Netherlands related to the utilization of AI for forecasting the likelihood of citizens committing social welfare fraud (*see* Politico article), this use of AI by governments has now also been categorized as "high-risk AI."

This session will discuss questions regarding the risk-based approach adopted by the proposed AI Act, imposed transparency obligations, the effectiveness of the fundamental rights impact assessment, and the ability to promote explainable and trustworthy AI.

Assigned reading

- Final compromise text of the proposed AI Act, recitals 31-44 (p. 19-36).
- Paul Friedl, & Gustavo Gil Gasiola, *Examining the EU's Artificial Intelligence Act*, (2024), https://verfassungsblog.de/examining-the-eus-artificial-intelligence-act/.
- Margot E. Kaminski, *Regulating the Risks of AI*, 103 B.U. L. Rev. 1347 (2023).
- Melissa Heikkilä, *Dutch Scandal Serves as a Warning for Europe over Risks of Using Algorithms*, Politico (2022), https://www.politico.eu/article/dutch-scandal-serves-as-a-warning-for-europe-over-risks-of-using-algorithms/.

Further reading

- Philipp Hacker, *AI Regulation in Europe: From the AI Act to Future Regulatory Challenges*, *in* Oxford handbook of algorithmic governance and the law (Ifeoma Ajunwa & Jeremias Adams-Prassl eds.).
- Johann Laux, Sandra Wachter, & Brent Mittelstadt, *Trustworthy Artificial Intelligence and the European Union AI Act: On the Conflation of Trustworthiness and Acceptability of Risk*, Regul. & Governance (2023).

IX. GOVERNANCE HOW? HUMAN IN THE LOOP

Human oversight is a component of different governance proposals. Human involvement in AI-enabled processes is also often thought to change something in our judgement about the quality or legitimacy of resulting decisions. This sessions dives into human-in-the-loop mechanisms to offer a better understanding of its role in governance proposals. It also offers more insight into what governance *by* algorithms means.

Assigned reading

- Rebecca Crootof, Margot E. Kaminski, & W. Nicholson Price, *Humans in the Loop*, 76 VAND. L. REV. 429 (2023).

Further reading

- Kiel Brennan-Marquez, Karen Levy, & Daniel Susser, *Strange Loops: Apparent Versus Actual Human Involvement in Automated Decision Making*, 34 BERKELEY TECH. L.J. 745 (2019).
- Deirdre K. Mulligan, & Helen Nissenbaum, *The Concept of Handoff as a Model for Ethical Analysis and Design*, *in* THE OXFORD HANDBOOK OF ETHICS OF AI 231 (Markus D. Dubber, Frank Pasquale, & Sunit Das eds., 2020).
- J. Nathan Matias, *Humans and Algorithms Work Together — so Study Them Together*, 617 NATURE 248 (2023).
- Ifeoma Ajunwa, *The Paradox of Automation as Anti-Bias Intervention*, 41 CARDOZO L. REV. 1671, 1707–1719 (2020).

X. GOVERNANCE HOW? ETHICS

Guest: Professor Ellen P. Goodman

This session considers ethics, both of professionals and as a field of study in the umbrella of AI. It will examine principles and codes of ethics and whether they stand for meaningful governance. It will also reflect on the role that professionals can play in the governance of AI, taking account of recent experiences with other professions.

Assigned reading

- Ellen P. Goodman, & Julia Trehu, *Algorithmic auditing: chasing AI accountability*, 39 SANTA CLARA HIGH TECH. L.J. 289, 296–306, 314–334 (2022) (Part I and Part III).
- Urs Gasser, & Carolyn Schmitt, *The Role of Professional Norms in the Governance of Artificial Intelligence*, *in* THE OXFORD HANDBOOK OF ETHICS OF AI 140 (Markus D. Dubber, Frank Pasquale, & Sunit Das eds., 2020).
- Luke Munn, *The Uselessness of AI Ethics*, AI AND ETHICS 1 (2022).
- John C. Coffee, *The Rise, Fall, and Redefinition of the Auditor: From Bookkeeper to Professional to Information Consultant*, *in* GATEKEEPERS: THE PROFESSIONS AND CORPORATE GOVERNANCE 108, 138–146, 161–171.

Further reading

- LOUISE AMOORE, CLOUD ETHICS: ALGORITHMS AND THE ATTRIBUTES OF OURSELVES AND OTHERS 85–107 (2020) (chapter 3).
- Nataliya Nedzhvetskaya, & J. S. Tan, *The Role of Workers in AI Ethics and Governance*, *in* THE OXFORD HANDBOOK OF AI GOVERNANCE C68.S1 (Justin B. Bullock et al. eds., 2024).
- Elettra Bietti, *From Ethics Washing to Ethics Bashing: A Moral Philosophy View on Tech Ethics*, 2 J. SOC. COMPUTING 266 (2021).
- Elizabeth Seger, *In Defence of Principlism in AI Ethics and Governance*, 35 PHIL. & TECH. 45 (2022).
- Jean-Christophe Bélisle-Pipon, et al. *Artificial Intelligence Ethics Has a Black Box Problem*, 38 AI & SOC'Y 1507 (2023).

XI. "TALKING AI LAW INTO BEING"

When we talk about the governance and regulation of AI, we are also talking about society and the law. How will they change in response to technology and the social, political, and legal orders made possible by it? This session looks inward and puts those questions on the foreground.

Assigned reading

- Aviv Ovadya, *Reimagining Democracy for AI*, 34 J. DEMOCRACY 162 (2023).
- Hannah Bloch-Wehba, *Algorithmic Governance from the Bottom up*, 48 BYU L. REV. 69, 109–136 (2022).
- Jack M. Balkin, *Free Speech Versus the First Amendment*, 70 UCLA L. REV. 1206, 44–51 (2023).
- Julie E. Cohen, BETWEEN TRUTH AND POWER 170–201 (2019).

Further reading

- Langdon Winner, *Do Artefacts Have Politics?*, *in* THE WHALE AND THE REACTOR: A SEARCH FOR LIMITS IN AN AGE OF HIGH TECHNOLOGY 19 (1986).
- Jascha Bareis, & Christian Katzenbach, *Talking AI into Being: The Narratives and Imaginaries of National AI Strategies and Their Performative Politics*, 47 SCI., TECH., & HUM. VALUES 855 (2022).
- Orly Lobel, *The Law of AI for Good*, SSRN ELECTRONIC JOURNAL (2023).
- James Grimmelmann, & Daniel Westreich, *Incomprehensible Discrimination*, 7 CAL. L. REV. ONLINE 164 (2017).
- Ngozi Okidegbe, *Discredited Data*, 107 CORNELL L. REV. 2007 (2022).

POLICIES & OTHER RELEVANT INFORMATION

DIVERSITY, EQUITY, INCLUSION, & BELONGING

As an instructor I am committed to making the class a community for equitable and inclusive learning. We will approach this collectively, as we owe it to each other. We can refer to the following statement by Dr. Rona Ramos, Lecturer and Graduate Services Coordinator at the Department of Physics.

> This class strives to be an inclusive community, learning from the many perspectives that come from having differing backgrounds and beliefs. As a community, we aim to be respectful to all. We reject all forms of prejudice and discrimination, including but not limited to those based on age, color, disability, gender, gender identity, gender expression, national origin, political affiliation, race, religion, sexual orientation, and veteran status. Faculty and students are expected to commit to creating an environment that facilitates inquiry and self-expression, while also demonstrating diligence in understanding how others' viewpoints may be different from their own.

ACCESSIBILITY

Equitable and inclusive teaching means accounting for accessibility for all students. We will all need accommodations because we all learn differently. If there are aspects of this course that prevent you from learning or that form barriers to your inclusion, please let me know as soon as possible. Together we'll develop strategies that can enable you to succeed in the course. I encourage you to visit Student Accessibility Services to determine how you could improve your learning as well. If you need official accommodations, you have a right to have these met. Undergraduate students can also benefit from the Poorvu Center's Academic Strategies Program resources for students with disabilities & neurodivergence.

ACADEMIC AND WRITING RESOURCES

Speaking from personal experience, I strongly encourage you to take advantage of the range of academic resources on campus, including options made available at the Poorvu Center (graduate and undergraduate students, as well as law-school specific offerings). Yale Jackson boasts its own Writing Program, headed by David Morse. Undergraduate students can also work with Residential College Tutors and benefit from the Academic Strategies Program. And, of course, I am also available to support you.

ACADEMIC INTEGRITY

Academic integrity is a core university value that ensures respect for the academic reputation of the University, its students, faculty and staff, and the degrees it confers. The University expects that students will conduct themselves in an honest and ethical manner and respect the intellectual work of others. Please ask about my expectations regarding permissible or encouraged forms of collaboration if they are unclear. Any work that you submit at any stage of the writing process — thesis, outline, draft, bibliography, final submission, presentations, blog posts, and more — must be your own. In addition, any words, ideas, or data that you borrow from other people and include in your work must be properly documented. Failure to do either of these things is plagiarism.

ACKNOWLEDGMENTS

For this syllabus I borrowed from related courses taught by Anupam Chander, Jennifer Cobbe, Anat Lior, Julian Posada, Ketan Ramakrishnan, and Andrew Selbst.

I am thankful to Dorrin Akbari, Victoria Hallinan, and Stephanie Ranks for feedback and support. The LaTeX template on which this syllabus was created is Kieran Healy's.